

```

1 !mkdir ~/.kaggle

1 !cp /content/kaggle.json ~/.kaggle/

1 !kaggle datasets download -d joosthazelzet/lego-brick-images

Warning: Your Kaggle API key is readable by other users on this system! To fix this, you can run 'chmod 600 /root/.kaggle/kaggle.json'
Downloading lego-brick-images.zip to /content
100% 1.00G/1.00G [00:38<00:00, 36.6MB/s]
100% 1.00G/1.00G [00:38<00:00, 27.7MB/s]

1 !unzip /content/lego-brick-images.zip

1 # %load_ext autoreload
2 # %autoreload 2
3 import numpy as np
4 import matplotlib.pyplot as plt
5
6 import tensorflow as tf
7 from tensorflow.keras import (
8     layers,
9     models,
10    callbacks,
11    losses,
12    utils,
13    metrics,
14    optimizers,
15 )
16
17 from drive.MyDrive.GAN_Datasets.utils import display, sample_batch

1 IMAGE_SIZE = 64
2 CHANNELS = 1
3 BATCH_SIZE = 128
4 Z_DIM = 100
5 EPOCHS = 300
6 LOAD_MODEL = False
7 ADAM_BETA_1 = 0.5
8 ADAM_BETA_2 = 0.999
9 LEARNING_RATE = 0.0002
10 NOISE_PARAM = 0.1

1 # Prepare the data
2 train_data = utils.image_dataset_from_directory(
3     "/content/LEGO brick images v1",
4     labels=None,
5     color_mode="grayscale",
6     image_size=(IMAGE_SIZE, IMAGE_SIZE),
7     batch_size=BATCH_SIZE,
8     shuffle=True,
9     seed=42,
10    interpolation="bilinear",
11 )

    Found 6379 files belonging to 1 classes.

1 def preprocess(img):
2     img = (tf.cast(img, "float32") - 127.5)/127.5
3     return img
4
5 train = train_data.map(lambda x:preprocess(x))

1 train_sample = sample_batch(train)
2 display(train_sample)

```



Build the GAN

```

1 discriminator_input = layers.Input(shape=(IMAGE_SIZE, IMAGE_SIZE, CHANNELS))
2 x = layers.Conv2D(64, kernel_size=4, strides=2, padding='same', use_bias=False)(discriminator_input)
3 x = layers.LeakyReLU(0.2)(x)
4 x = layers.Dropout(0.3)(x)
5 x = layers.Conv2D(128, kernel_size=4, strides=2, padding='same', use_bias=False)(x)
6 x = layers.BatchNormalization(momentum=0.9)(x)
7 x = layers.LeakyReLU(0.2)(x)
8 x = layers.Dropout(0.3)(x)
9 x = layers.Conv2D(256, kernel_size=4, strides=2, padding='same', use_bias=False)(x)
10 x = layers.BatchNormalization(momentum=0.9)(x)
11 x = layers.LeakyReLU(0.2)(x)
12 x = layers.Dropout(0.3)(x)
13 x = layers.Conv2D(512, kernel_size=4, strides=2, padding='same', use_bias=False)(x)
14 x = layers.BatchNormalization(momentum=0.9)(x)
15 x = layers.LeakyReLU(0.2)(x)
16 x = layers.Dropout(0.3)(x)
17 x = layers.Conv2D(1, kernel_size=4, strides=1, padding='valid', use_bias=False, activation='sigmoid')(x)
18 discriminator_output = layers.Flatten()(x)
19 discriminator = models.Model(discriminator_input, discriminator_output)
20
21 discriminator.summary()

```

Model: "model"

Layer (type)	Output Shape	Param #
=====		
input_1 (InputLayer)	[(None, 64, 64, 1)]	0
conv2d (Conv2D)	(None, 32, 32, 64)	1024
leaky_re_lu (LeakyReLU)	(None, 32, 32, 64)	0
dropout (Dropout)	(None, 32, 32, 64)	0
conv2d_1 (Conv2D)	(None, 16, 16, 128)	131072
batch_normalization (Batch Normalization)	(None, 16, 16, 128)	512
leaky_re_lu_1 (LeakyReLU)	(None, 16, 16, 128)	0
dropout_1 (Dropout)	(None, 16, 16, 128)	0
conv2d_2 (Conv2D)	(None, 8, 8, 256)	524288
batch_normalization_1 (Batch Normalization)	(None, 8, 8, 256)	1024
leaky_re_lu_2 (LeakyReLU)	(None, 8, 8, 256)	0
dropout_2 (Dropout)	(None, 8, 8, 256)	0
conv2d_3 (Conv2D)	(None, 4, 4, 512)	2097152
batch_normalization_2 (Batch Normalization)	(None, 4, 4, 512)	2048
leaky_re_lu_3 (LeakyReLU)	(None, 4, 4, 512)	0
dropout_3 (Dropout)	(None, 4, 4, 512)	0
conv2d_4 (Conv2D)	(None, 1, 1, 1)	8192
flatten (Flatten)	(None, 1)	0

=====

Total params: 2,765,312

Trainable params: 2,763,520

Non-trainable params: 1,792

```

1 generator_input = layers.Input(shape=(Z_DIM,))
2 x = layers.Reshape((1,1,Z_DIM))(generator_input)
3 x = layers.Conv2DTranspose(512, kernel_size=4, strides=1, padding='valid', use_bias=False)(x)
4 x = layers.BatchNormalization(momentum=0.9)(x)
5 x = layers.LeakyReLU(0.2)(x)
6 x = layers.Conv2DTranspose(256, kernel_size=4, strides=2, padding='same', use_bias=False)(x)
7 x = layers.BatchNormalization(momentum=0.9)(x)
8 x = layers.LeakyReLU(0.2)(x)
9 x = layers.Conv2DTranspose(128, kernel_size=4, strides=2, padding='same', use_bias=False)(x)
10 x = layers.BatchNormalization(momentum=0.9)(x)
11 x = layers.LeakyReLU(0.2)(x)
12 x = layers.Conv2DTranspose(64, kernel_size=4, strides=2, padding='same', use_bias=False)(x)
13 x = layers.BatchNormalization(momentum=0.9)(x)
14 x = layers.LeakyReLU(0.2)(x)
15 generator_output = layers.Conv2DTranspose(CHANNELS, kernel_size=4, strides=2, padding='same', use_bias=False, activation='tanh')(x)
16 generator = models.Model(generator_input, generator_output)
17 generator.summary()

```

Model: "model_1"

Layer (type)	Output Shape	Param #
=====		
input_2 (InputLayer)	[(None, 100)]	0
reshape (Reshape)	(None, 1, 1, 100)	0
conv2d_transpose (Conv2DTra nspose)	(None, 4, 4, 512)	819200
batch_normalization_3 (Batc hNormalization)	(None, 4, 4, 512)	2048
leaky_re_lu_4 (LeakyReLU)	(None, 4, 4, 512)	0
conv2d_transpose_1 (Conv2DT ranspose)	(None, 8, 8, 256)	2097152
batch_normalization_4 (Batc hNormalization)	(None, 8, 8, 256)	1024
leaky_re_lu_5 (LeakyReLU)	(None, 8, 8, 256)	0
conv2d_transpose_2 (Conv2DT ranspose)	(None, 16, 16, 128)	524288
batch_normalization_5 (Batc hNormalization)	(None, 16, 16, 128)	512
leaky_re_lu_6 (LeakyReLU)	(None, 16, 16, 128)	0
conv2d_transpose_3 (Conv2DT ranspose)	(None, 32, 32, 64)	131072
batch_normalization_6 (Batc hNormalization)	(None, 32, 32, 64)	256
leaky_re_lu_7 (LeakyReLU)	(None, 32, 32, 64)	0
conv2d_transpose_4 (Conv2DT ranspose)	(None, 64, 64, 1)	1024

```

=====
Total params: 3,576,576
Trainable params: 3,574,656
Non-trainable params: 1,920

```

```

1 class DCGAN(models.Model):
2     def __init__(self, discriminator, generator, latent_dim):
3         super(DCGAN, self).__init__()
4         self.discriminator = discriminator
5         self.generator = generator
6         self.latent_dim = latent_dim
7
8     def compile(self, d_optimizer, g_optimizer):
9         super(DCGAN, self).compile()
10        self.loss_fn = losses.BinaryCrossentropy()
11        self.d_optimizer = d_optimizer
12        self.g_optimizer = g_optimizer
13        self.d_loss_metric = metrics.Mean(name="d_loss")
14        self.d_real_acc_metric = metrics.BinaryAccuracy(name="d_real_acc")
15        self.d_fake_acc_metric = metrics.BinaryAccuracy(name="d_fake_acc")
16        self.d_acc_metric = metrics.BinaryAccuracy(name="d_acc")
17        self.g_loss_metric = metrics.Mean(name="g_loss")
18        self.g_acc_metric = metrics.BinaryAccuracy(name="g_acc")
19
20    @property
21    def metrics(self):
22        return [
23            self.d_loss_metric,
24            self.d_real_acc_metric,
25            self.d_fake_acc_metric,
26            self.d_acc_metric,
27            self.g_loss_metric,
28            self.g_acc_metric,
29        ]
30
31    def train_step(self, real_images):
32        # Sample random points in the latent space
33        batch_size = tf.shape(real_images)[0]
34        random_latent_vectors = tf.random.normal(
35            shape=(batch_size, self.latent_dim)
36        )
37
38        # Train the discriminator on fake images
39        with tf.GradientTape() as gen_tape, tf.GradientTape() as disc_tape:
40            generated_images = self.generator(
41                random_latent_vectors, training=True
42            )
43            real_predictions = self.discriminator(real_images, training=True)
44            fake_predictions = self.discriminator(
45                generated_images, training=True
46            )
47
48            real_labels = tf.ones_like(real_predictions)
49            real_noisy_labels = real_labels + NOISE_PARAM * tf.random.uniform(
50                tf.shape(real_predictions)
51            )
52            fake_labels = tf.zeros_like(fake_predictions)
53            fake_noisy_labels = fake_labels - NOISE_PARAM * tf.random.uniform(
54                tf.shape(fake_predictions)
55            )
56
57            d_real_loss = self.loss_fn(real_noisy_labels, real_predictions)
58            d_fake_loss = self.loss_fn(fake_noisy_labels, fake_predictions)
59            d_loss = (d_real_loss + d_fake_loss) / 2.0
60
61            g_loss = self.loss_fn(real_labels, fake_predictions)
62
63            gradients_of_discriminator = disc_tape.gradient(
64                d_loss, self.discriminator.trainable_variables
65            )
66            gradients_of_generator = gen_tape.gradient(
67                g_loss, self.generator.trainable_variables
68            )
69
70            self.d_optimizer.apply_gradients(
71                zip(gradients_of_discriminator, discriminator.trainable_variables)
72            )
73            self.g_optimizer.apply_gradients(
74                zip(gradients_of_generator, generator.trainable_variables)
75            )
76
77        # Update metrics

```

```

78         self.d_loss_metric.update_state(d_loss)
79         self.d_real_acc_metric.update_state(real_labels, real_predictions)
80         self.d_fake_acc_metric.update_state(fake_labels, fake_predictions)
81         self.d_acc_metric.update_state(
82             [real_labels, fake_labels], [real_predictions, fake_predictions]
83         )
84         self.g_loss_metric.update_state(g_loss)
85         self.g_acc_metric.update_state(real_labels, fake_predictions)
86
87         return {m.name: m.result() for m in self.metrics}

```

```

1 # create a DCGAN
2 dcgan = DCGAN(discriminator=discriminator, generator=generator , latent_dim=Z_DIM)

1 if LOAD_MODEL:
2     dcgan.load_weights("./checkpoint/checkpoint.ckpt")

```

Train the GAN

```


1 dcgan.compile(
2     d_optimizer=optimizers.Adam(
3         learning_rate=LEARNING_RATE, beta_1=ADAM_BETA_1, beta_2=ADAM_BETA_2
4     ),
5     g_optimizer=optimizers.Adam(
6         learning_rate=LEARNING_RATE, beta_1=ADAM_BETA_1, beta_2=ADAM_BETA_2
7     ),
8 )

1 model_checkpoint_callback = callbacks.ModelCheckpoint(
2     filepath="./checkpoint/checkpoint.ckpt",
3     save_weights_only=True,
4     save_freq="epoch",
5     verbose=0,
6 )
7
8 tensorboard_callback = callbacks.TensorBoard(log_dir="./logs")
9
10 class ImageGenerator(callbacks.Callback):
11     def __init__(self, num_img, latent_dim):
12         self.num_img = num_img
13         self.latent_dim = latent_dim
14
15     def on_epoch_end(self, epoch, logs=None):
16         random_latent_vectors = tf.random.normal(
17             shape=(self.num_img, self.latent_dim)
18         )
19         generated_images = self.model.generator(random_latent_vectors)
20         generated_images = generated_images * 127.5 + 127.5
21         generated_images = generated_images.numpy()
22         display(
23             generated_images,
24             save_to="/content/output/generated_img_%03d.png" % (epoch),
25         )

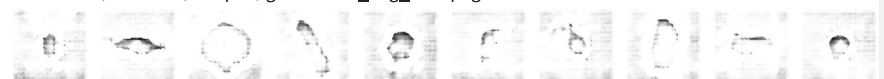
1 dcgan.fit(
2     train,
3     epochs=EPOCHS,
4     callbacks=[
5         model_checkpoint_callback,
6         tensorboard_callback,
7         ImageGenerator(num_img=10, latent_dim=Z_DIM),
8     ],
9 )

```

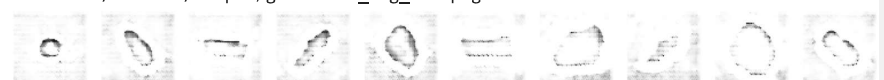
50/50 [=====] - 12s 210ms/step - d_loss: -0.7727 - d_real_ac
Epoch 286/300
50/50 [=====] - ETA: 0s - d_loss: -0.7655 - d_real_acc: 1.00
Saved to /content/output/generated_img_285.png




50/50 [=====] - 12s 211ms/step - d_loss: -0.7655 - d_real_ac
Epoch 287/300
50/50 [=====] - ETA: 0s - d_loss: -0.7652 - d_real_acc: 1.00
Saved to /content/output/generated_img_286.png



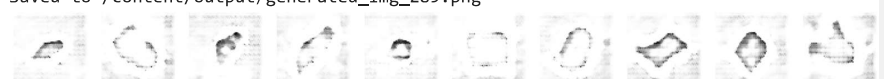
50/50 [=====] - 12s 211ms/step - d_loss: -0.7652 - d_real_ac
Epoch 288/300
50/50 [=====] - ETA: 0s - d_loss: -0.7698 - d_real_acc: 1.00
Saved to /content/output/generated_img_287.png



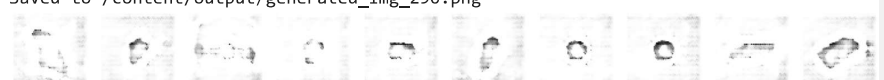
50/50 [=====] - 12s 211ms/step - d_loss: -0.7698 - d_real_ac
Epoch 289/300
50/50 [=====] - ETA: 0s - d_loss: -0.7651 - d_real_acc: 1.00
Saved to /content/output/generated_img_288.png



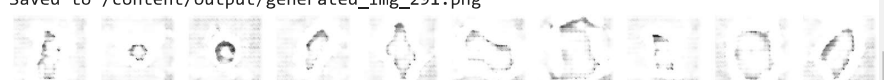
50/50 [=====] - 12s 216ms/step - d_loss: -0.7651 - d_real_ac
Epoch 290/300
50/50 [=====] - ETA: 0s - d_loss: -0.7713 - d_real_acc: 1.00
Saved to /content/output/generated_img_289.png



50/50 [=====] - 12s 209ms/step - d_loss: -0.7713 - d_real_ac
Epoch 291/300
50/50 [=====] - ETA: 0s - d_loss: -0.7660 - d_real_acc: 1.00
Saved to /content/output/generated_img_290.png



50/50 [=====] - 12s 221ms/step - d_loss: -0.7660 - d_real_ac
Epoch 292/300
50/50 [=====] - ETA: 0s - d_loss: -0.7726 - d_real_acc: 1.00
Saved to /content/output/generated_img_291.png



50/50 [=====] - 12s 214ms/step - d_loss: -0.7726 - d_real_ac
Epoch 293/300